



Claudio Mirabello

Intro to Glorified Function Fitting (a.k.a. Neural Networks and Deep Learning)



Resources


- ▶ MIT lectures on Deep Learning
(<http://introtodeeplearning.com/>)
- ▶ TensorFlow Playground
(<https://playground.tensorflow.org>)
- ▶ Keras Docs
(<https://keras.io>)




A quick note about what we're going to do

- ▶ We're going to cover mostly classification because it is easy to visualize and the most common task
- ▶ Regression in the end is not that different, as we will see it is just about changing one line of code
- ▶ (and your dataset)
- ▶ (and target function)
- ▶ (and your dataset)
- ▶ (but it is not that hard, I promise)


This is a neuron



This is a perceptron (1958)



Mark I Perceptron machine
wikipedia.org



Output

Linear combination of inputs


$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

Non-linear activation function


Bias




Perceptrons can only learn linearly separable classes




Minsky and Papert, 1969





Perceptrons can only learn linearly separable classes





But sometimes you want
to model non-linear functions





Ok, let's make the classifier
non-linear then?
(duh!)



Non-linear activation functions

Activation Functions

$$\hat{y} = g\left(w_0 + \mathbf{X}^T \mathbf{W}\right)$$

- Example: sigmoid function


$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$





Common activation functions


Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$


$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$




Special case: softmax

- Used in classification problems
- Given k classes, it decides which one is more likely
- One output per class, each output is assigned a probability from 0 to 1
- The sum of probabilities for all outputs is 1

$$\sigma(\mathbf{z})_j = \frac{e^{z_j}}{\sum_{k=1}^k e^{z_k}} \text{ for } j = 1, \dots, k$$




Wait a second, the perceptron already has
a non-linear activation function!



That was in the 50s

- ▶ Perceptron can't deal with many types of more complex patterns (XOR)
- ▶ Now we have huge, complex datasets (Big Data)
- ▶ Faster processors, more memory
- ▶ And a better understanding of learning processes

This is a perceptron




Output

Linear combination of inputs


$$\hat{y} = g \left(w_0 + \sum_{i=1}^m x_i w_i \right)$$

activation function

Bias



Multi-layer Perceptron (1986)




Inputs


Hidden


Final Output

$$z_i = w_{0,i}^{(1)} + \sum_{j=1}^m x_j w_{j,i}^{(1)} \quad \hat{y}_i = g \left(w_{0,i}^{(2)} + \sum_{j=1}^{d_1} z_j w_{j,i}^{(2)} \right)$$




Now we're getting somewhere



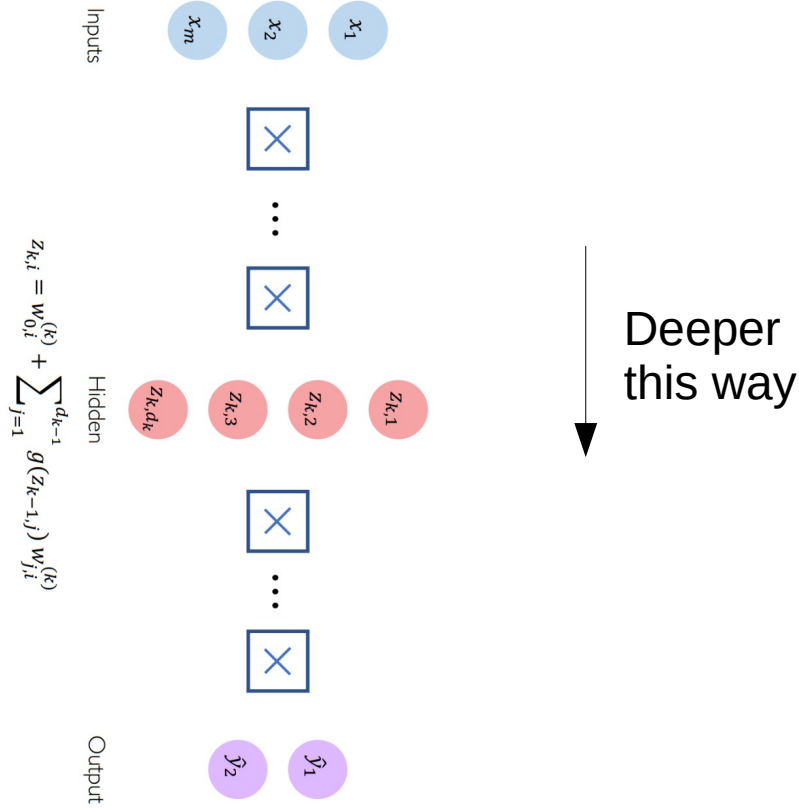


Why stop at one hidden layer?




$$z_{k,i} = w_{0,i}^{(k)} + \sum_{j=1}^{d_{k-1}} g(z_{k-1,j}) w_{j,i}^{(k)}$$

Deep Networks are simply NNs with multiple hidden layers




<https://playground.tensorflow.org>

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.



Let's review:

- Perceptron
- XOR problem
- Activations
- Multi-layer perceptron




How do we decide which weights are optimal?

- A linear regressor's weights (coefficients) are calculated in closed form
- This can't be done if you have hidden layers and non-linear activations

How do we decide which weights are optimal?

The **loss** of our network measures the cost incurred from incorrect predictions




$$\mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{W})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$



How do we decide which weights are optimal?


The **empirical loss** measures the total loss over our entire dataset



Also known as:

- Objective function
- Cost function
- Empirical Risk

$$J(\mathbf{W}) = \frac{1}{n} \sum_{i=1}^n \mathcal{L}(\underbrace{f(x^{(i)}; \mathbf{W})}_{\text{Predicted}}, \underbrace{y^{(i)}}_{\text{Actual}})$$




Lower loss => better predictions

We want to find the network weights that **achieve the lowest loss**

$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} \frac{1}{n} \sum_{i=1}^n \mathcal{L}(f(x^{(i)}; \mathbf{W}), y^{(i)})$$


$$\mathbf{W}^* = \operatorname{argmin}_{\mathbf{W}} J(\mathbf{W})$$




Minimizing loss

$$\mathbf{W}^* = \underset{\mathbf{W}}{\operatorname{argmin}} J(\mathbf{W})$$


Remember:
*Our loss is a function of
the network weights!*






Minimizing loss


Randomly pick an initial (w_0, w_1)






Minimizing loss


Compute gradient, $\frac{\partial J(W)}{\partial W}$






Minimizing loss


Take small step in opposite direction of gradient





Minimizing loss

Repeat until convergence






Gradient descent


Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights



Activation functions have to be differentiable!


Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


Hyperbolic Tangent



$$g(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$


$$g'(z) = 1 - g(z)^2$$

Rectified Linear Unit (ReLU)



$$g(z) = \max(0, z)$$

$$g'(z) = \begin{cases} 1, & z > 0 \\ 0, & \text{otherwise} \end{cases}$$





The learning rate η

Algorithm


1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
4. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
5. Return weights

Take small step in opposite direction of gradient







Backpropagation



$$\frac{\partial J(W)}{\partial w_2} = \underbrace{\frac{\partial J(W)}{\partial \hat{y}}}_{\text{purple}} * \underbrace{\frac{\partial \hat{y}}{\partial w_2}}_{\text{red}}$$



Backpropagation




$$\frac{\partial J(\mathbf{W})}{\partial w_1} = \underbrace{\frac{\partial J(\mathbf{W})}{\partial \hat{y}}}_{\text{purple}} * \underbrace{\frac{\partial \hat{y}}{\partial z_1}}_{\text{red}} * \underbrace{\frac{\partial z_1}{\partial w_1}}_{\text{blue}}$$

Backpropagation example, step by step:

<https://mattmazur.com/2015/03/17/a-step-by-step-backpropagation-example/>

1. Forward pass



$$\text{net}_{h1} = 0.05 * 0.14 + 0.1 * 0.19$$

$$\text{out}_{h1} = 1/(1+e^{-\text{net}_{h1}})$$


$$\text{net}_{h2} = 0.05 * 0.24 + 0.1 * 0.29$$

$$\text{out}_{h2} = 1/(1+e^{-\text{net}_{h2}})$$

$$\text{net}_{o1} = 0.35 * \text{h1} + 0.4 * \text{h2}$$

$$\text{out}_{o1} = 1/(1+e^{-\text{net}_{o1}})$$

1. Forward pass



$$\text{net}_{h1} = 0.05 * 0.14 + 0.1 * 0.19 = 0.026$$

$$\text{out}_{h1} = 1/(1+e^{-\text{net}_{h1}}) = 0.50$$


$$\text{net}_{h2} = 0.05 * 0.24 + 0.1 * 0.29 = 0.041$$

$$\text{out}_{h2} = 1/(1+e^{-\text{net}_{h2}}) = 0.51$$

$$\text{net}_{o1} = 0.35 * h1 + 0.4 * h2 = 0.38$$

$$\text{out}_{o1} = 1/(1+e^{-\text{net}_{o1}}) = 0.59$$

2. Calculate error



$$\text{net}_{h1} = 0.05 * 0.14 + 0.1 * 0.19 = 0.026$$

$$\text{out}_{h1} = 1/(1+e^{-\text{net}_{h1}}) = 0.50$$

$$\text{net}_{h2} = 0.05 * 0.24 + 0.1 * 0.29 = 0.041$$


$$\text{out}_{h2} = 1/(1+e^{-\text{net}_{h2}}) = 0.51$$

$$\text{net}_{o1} = 0.35 * h1 + 0.4 * h2 = 0.38$$

$$\text{out}_{o1} = 1/(1+e^{-\text{net}_{o1}}) = 0.59$$


$$E_{o1} = \frac{1}{2} * (0.1 - 0.59)^2 = 0.12$$

3. Derivatives



$$\begin{aligned} \frac{\partial E_{o1}}{\partial w5} &= \frac{\partial E_{o1}}{\partial \text{out}_{o1}} && \text{(error)} \\ &\quad * \frac{\partial o1}{\partial \text{net}_{o1}} && \text{(activation)} \\ &\quad * \frac{\partial \text{net}_{o1}}{\partial w5} && \text{(weight)} \end{aligned}$$


3. Derivatives



$$\begin{aligned} \frac{\partial E_{o1}}{\partial w5} &= \frac{\partial E_{o1}}{\partial out_{o1}} \quad \text{(error)} \\ &\quad * \frac{\partial o1}{\partial net_{o1}} \quad \text{(activation)} \\ &\quad * \frac{\partial net_{o1}}{\partial w5} \quad \text{(weight)} \end{aligned}$$

$$\begin{aligned} E_{o1} &= \frac{1}{2} * (o1 - \hat{o}1)^2 \\ \frac{\partial E_{o1}}{\partial out_{o1}} &= 2 * \frac{1}{2} * (o1 - \hat{o}1) * -1 \\ &= -o1 + \hat{o}1 = -0.1 + 0.59 = 0.49 \end{aligned}$$


3. Derivatives




$$\begin{aligned} \partial E_{o1} / \partial w5 &= 0.49 \\ &\quad * \partial o1 / \partial \text{net}_{o1} \quad \text{(error)} \\ &\quad * \partial \text{net}_{o1} / \partial w5 \quad \text{(activation)} \end{aligned}$$

$$\begin{aligned} E_{o1} &= \frac{1}{2} * (o1 - \hat{o}1)^2 \\ \partial E_{o1} / \partial \text{out}_{o1} &= 2 * \frac{1}{2} * (o1 - \hat{o}1) * -1 \\ &= -o1 + \hat{o}1 = -0.1 + 0.59 = 0.49 \end{aligned}$$

3. Derivatives



Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


$$\begin{aligned} \frac{\partial E_{o1}}{\partial w5} &= 0.49 && \text{(error)} \\ &\times \frac{\partial o1}{\partial \text{net}_{o1}} && \text{(activation)} \\ &\times \frac{\partial \text{net}_{o1}}{\partial w5} && \text{(weight)} \end{aligned}$$

$$\begin{aligned} o1 &= 1/(1+e^{-\text{net}_{o1}}) \\ \frac{\partial o1}{\partial \text{net}_{o1}} &= 1/(1+e^{-\text{net}_{o1}}) * (1 - 1/(1+e^{-\text{net}_{o1}})) \\ &= 0.24 \end{aligned}$$

3. Derivatives



Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


$$\begin{aligned} \frac{\partial E_{o1}}{\partial w5} &= 0.49 && \text{(error)} \\ &\times 0.24 && \text{(activation)} \\ &\times \frac{\partial \text{net}_{o1}}{\partial w5} && \text{(weight)} \end{aligned}$$

$$\begin{aligned} o1 &= 1/(1+e^{-\text{net}_{o1}}) \\ \frac{\partial o1}{\partial \text{net}_{o1}} &= 1/(1+e^{-\text{net}_{o1}}) * (1 - 1/(1+e^{-\text{net}_{o1}})) \\ &= 0.24 \end{aligned}$$

3. Derivatives



Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


$$\begin{aligned} \partial E_{o1} / \partial w5 &= 0.49 && \text{(error)} \\ &* 0.24 && \text{(activation)} \\ &* \partial \text{net}_{o1} / \partial w5 && \text{(weight)} \end{aligned}$$

$$\begin{aligned} \text{net}_{o1} &= w5 * \text{out}_{h1} \\ \partial \text{net}_{o1} / \partial w5 &= w5 = 0.35 \end{aligned}$$

3. Derivatives



Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$


$$g'(z) = g(z)(1 - g(z))$$

$$\begin{aligned} \partial E_{o1} / \partial w5 &= 0.49 \\ &\quad * 0.24 \\ &\quad * 0.35 \end{aligned}$$


(error)
(activation)
(weight)

$$\begin{aligned} \text{net}_{o1} &= w5 * \text{out}_{h1} \\ \partial \text{net}_{o1} / \partial w5 &= w5 = 0.35 \end{aligned}$$

3. Derivatives



Sigmoid Function




$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


$$\begin{aligned}\partial E_{o1} / \partial w5 &= 0.49 \\ &\quad * 0.24 \\ &\quad * 0.35 \\ &= 0.04\end{aligned}$$

(error)
(activation)
(weight)
(w5 gradient)

4. Weight update



Sigmoid Function



$$g(z) = \frac{1}{1 + e^{-z}}$$

$$g'(z) = g(z)(1 - g(z))$$


$$\begin{aligned} \partial E_{o1} / \partial w5 &= 0.49 \\ &\quad * 0.24 \\ &\quad * 0.35 \\ &= 0.04 \end{aligned}$$


(error)
(activation)
(weight)
(w5 gradient)

$$w5' = w5 - \eta * 0.04 = 0.35 - 0.04 = 0.31$$

Exercise

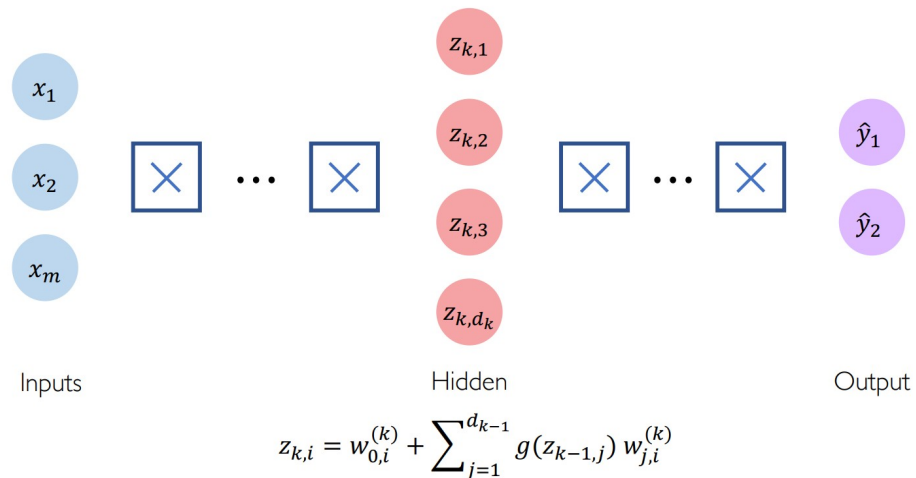
- ▶ Can you calculate the weight update for w_6 ? How many new gradients do you need to calculate?
- ▶ Can you go one step back and calculate the weight update for w_1/w_4 ?
- ▶ What is the new predicted output? Has the error gone down?






Gradient vanishing

- ▶ What happens if we backpropagate on a network with many hidden layers?



$$\partial E_{y_1} / \partial w_1 = \partial E_{o_1} / \partial o_{o_1} * \partial o_1 / \partial \text{net}_{o_1} * \dots * \partial E_{o_1} / \partial o_{o_1} * \partial o_1 / \partial \text{net}_{o_1} * \partial \text{net}_{o_1} / \partial w_1$$



Gradient vanishing


- ▶ These are all “zero-point-somethings” multiplied by each other
- ▶ So the gradient becomes smaller by orders of magnitudes as we go back more and more layers until it's so small that the network is stuck

$$\partial E_{y1} / \partial w_1 = \partial E_{o1} / \partial o_{o1} * \partial o_1 / \partial net_{o1} * \dots * \partial E_{o1} / \partial o_{o1} * \partial o_1 / \partial net_{o1} * \partial net_{o1} / w_1 = 0.0000000001$$

initial $w_1 = 0.5$


optimal $w_1 = -0.2$

How many iterations do we need to get from 0.5 to -0.2?





Other issues

- ▶ Finding the right learning rate (η , *eta*) can be tricky
- ▶ Computing the loss over the full set at every iteration is slow





Too small: slow, gets stuck







Too large: fast, might explode





Adaptive learning rate







Stochastic gradient descent


Algorithm

1. Initialize weights randomly $\sim \mathcal{N}(0, \sigma^2)$
2. Loop until convergence:
3. Pick batch of B data points
4. Compute gradient, $\frac{\partial J(\mathbf{W})}{\partial \mathbf{W}} = \frac{1}{B} \sum_{k=1}^B \frac{\partial J_k(\mathbf{W})}{\partial \mathbf{W}}$
5. Update weights, $\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial J(\mathbf{W})}{\partial \mathbf{W}}$
6. Return weights






Other issues: under/overfitting




Underfitting

Model does not have capacity to fully learn the data




Ideal fit



Overfitting

Too complex, extra parameters, does not generalize well



Regularization to limit overfitting

L1 (Lasso):


$$L(x, y) \equiv \sum_{i=1}^n (y_i - h_{\theta}(x_i))^2 + \lambda \sum_{i=1}^n |\theta_i|$$

L2 (Ridge):


$$L(x, y) \equiv \sum_{i=1}^n (y_i - h_{\theta}(x_i))^2 + \lambda \sum_{i=1}^n \theta_i^2$$

Regularization to limit overfitting

Early stopping




Dropout




<https://playground.tensorflow.org>

Tinker With a **Neural Network** Right Here in Your Browser.
Don't Worry, You Can't Break It. We Promise.




Let's review:

- Learning rate
- Minibatches
- Regularization
- Under/Overfitting



Before we continue

- ▶ Go to: https://bitbucket.org/clami66/deep_learning_course
- ▶ Clone the repository with git
(git clone [git@bitbucket.org](https://bitbucket.org):clami66/deep_learning_course.git
git clone
https://clami66@bitbucket.org/clami66/deep_learning_course.git)
- ▶ Or use the “Download Repository” option
(left menu → Downloads → Download Repository)
- ▶ You will get a Jupyter Notebook as well as the latest version of the slides



Visual cortex studies ('50s)

- ▶ Hubel & Wiesel (1959) hypothesized that the visual cortex is made of two types of cells:
 - ▶ Simple cells recognize fixed patterns such as edges
 - ▶ Complex cells can do the same task, but in a spatially invariant fashion
 - ▶ A model was proposed where simple and complex cells were combined in a cascade

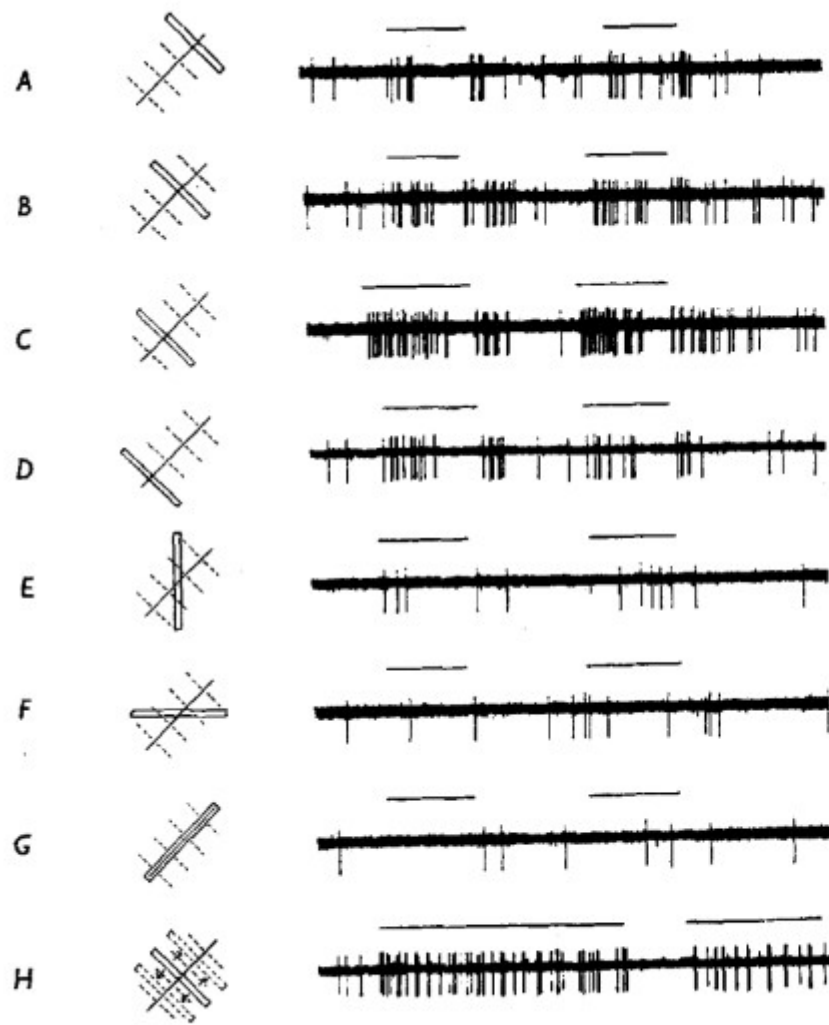




Fig. 4. Responses of a cell with a complex field to stimulation of the left (contralateral) eye with a slit $\frac{1}{8} \times 2\frac{1}{2}^\circ$. Receptive field was in the area centralis and was about $2 \times 3^\circ$ in size. A–D, $\frac{1}{8}^\circ$ wide slit oriented parallel to receptive field axis. E–G, slit oriented at 45 and 90° to receptive-field axis. H, slit oriented as in A–D, is on throughout the record and is moved rapidly from side to side where indicated by upper beam. Responses from left eye slightly more marked than those from right (Group 3, see Part II). Time 1 sec.



Neocognitron (1979)






Image processing: kernels

Input image



Convolution
Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map







Image processing: kernels








Image processing: kernels






Convnets apply adaptive filters to the inputs



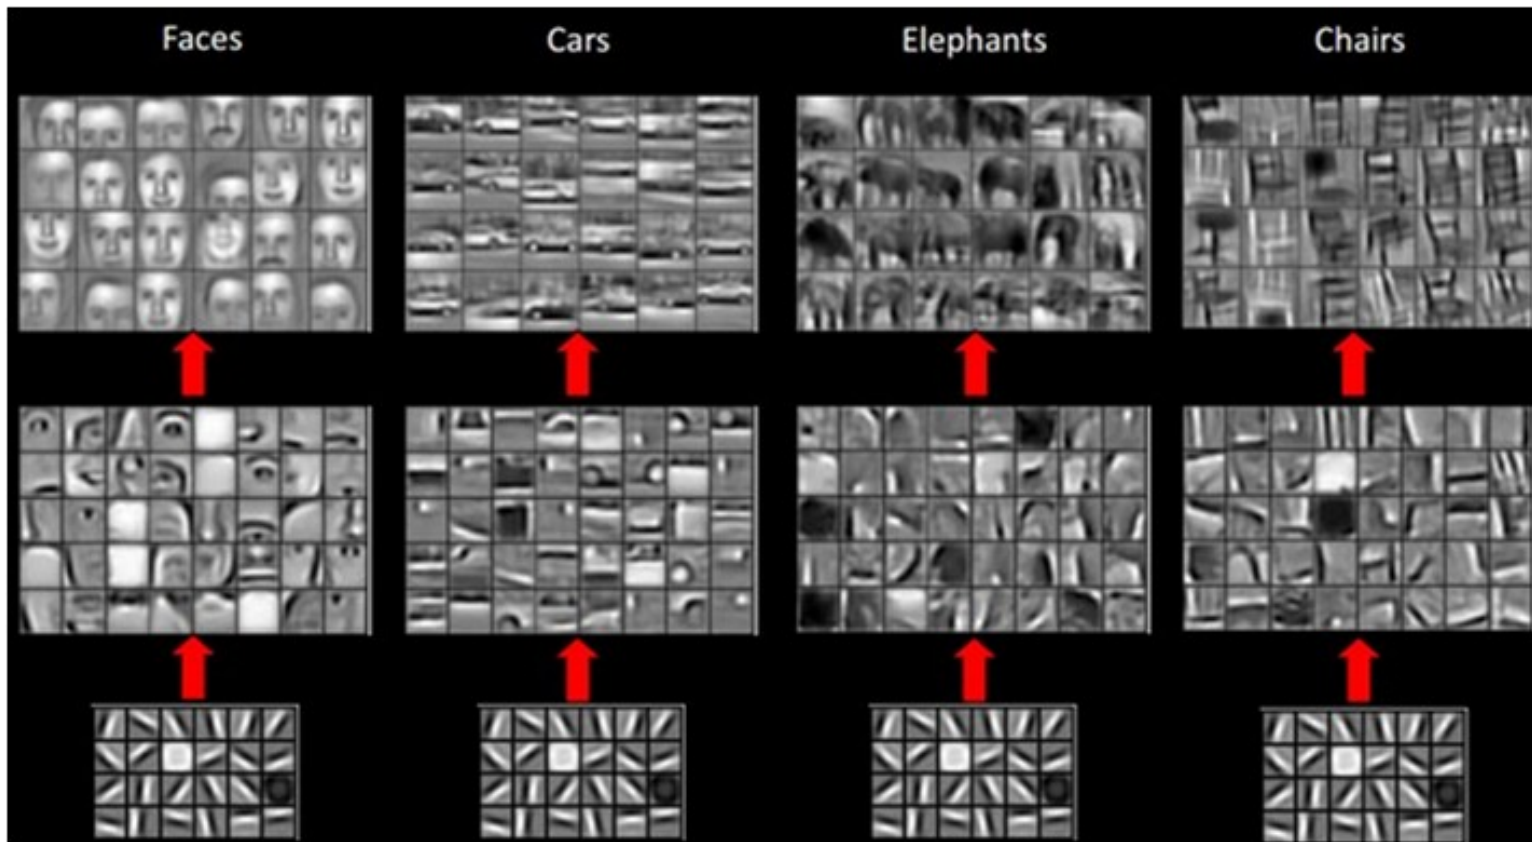


Convolutional Neural Networks (1995)



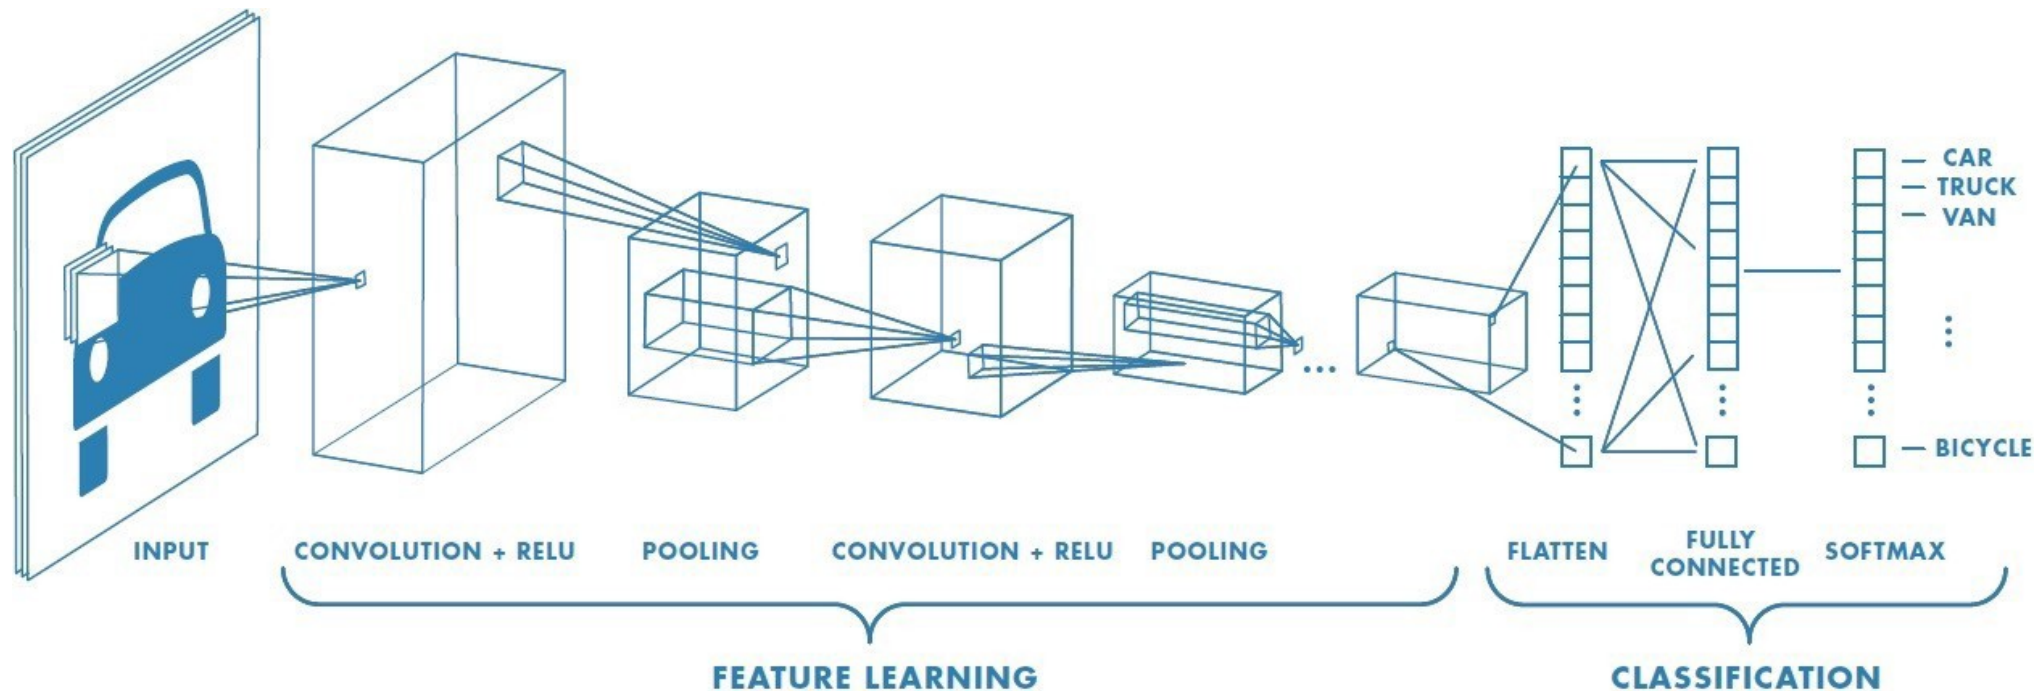


Convnets apply adaptive filters to inputs

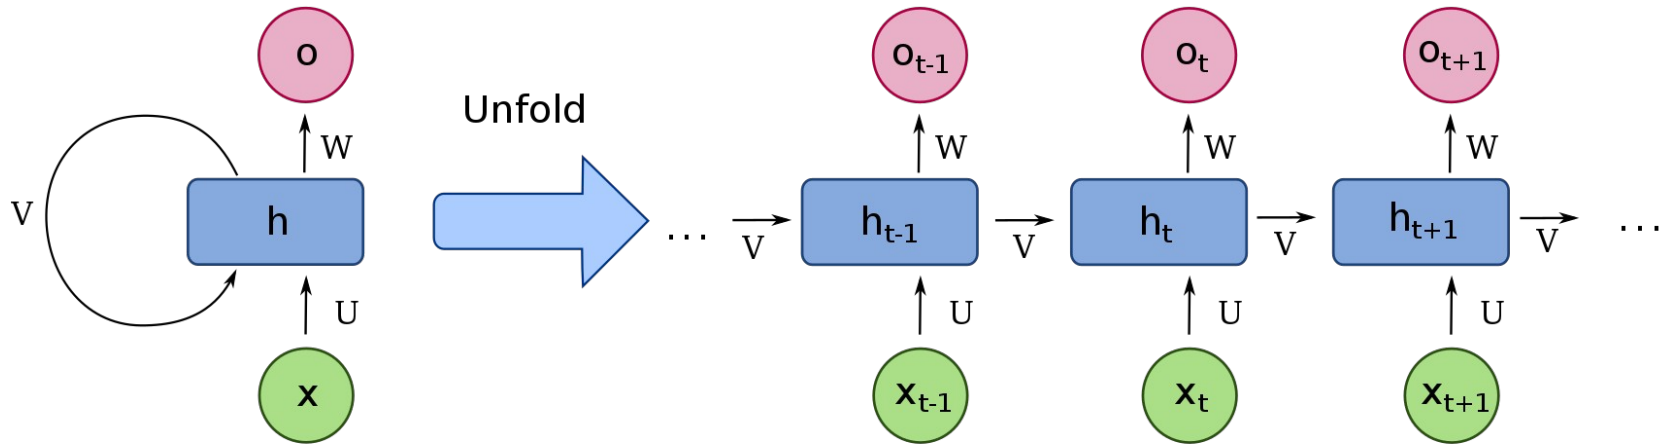





Convolutional Neural Networks (1995)




Recurrent Neural Networks




- Connections between nodes form a directed graph along a temporal sequence
- Actually always the same neurons repeated over and over (weight sharing)
- RNNs work very well when your data is like a time series (audio, text, DNA bases...)




Word embeddings (word2vec)



Male-Female




Verb tense



Country-Capital

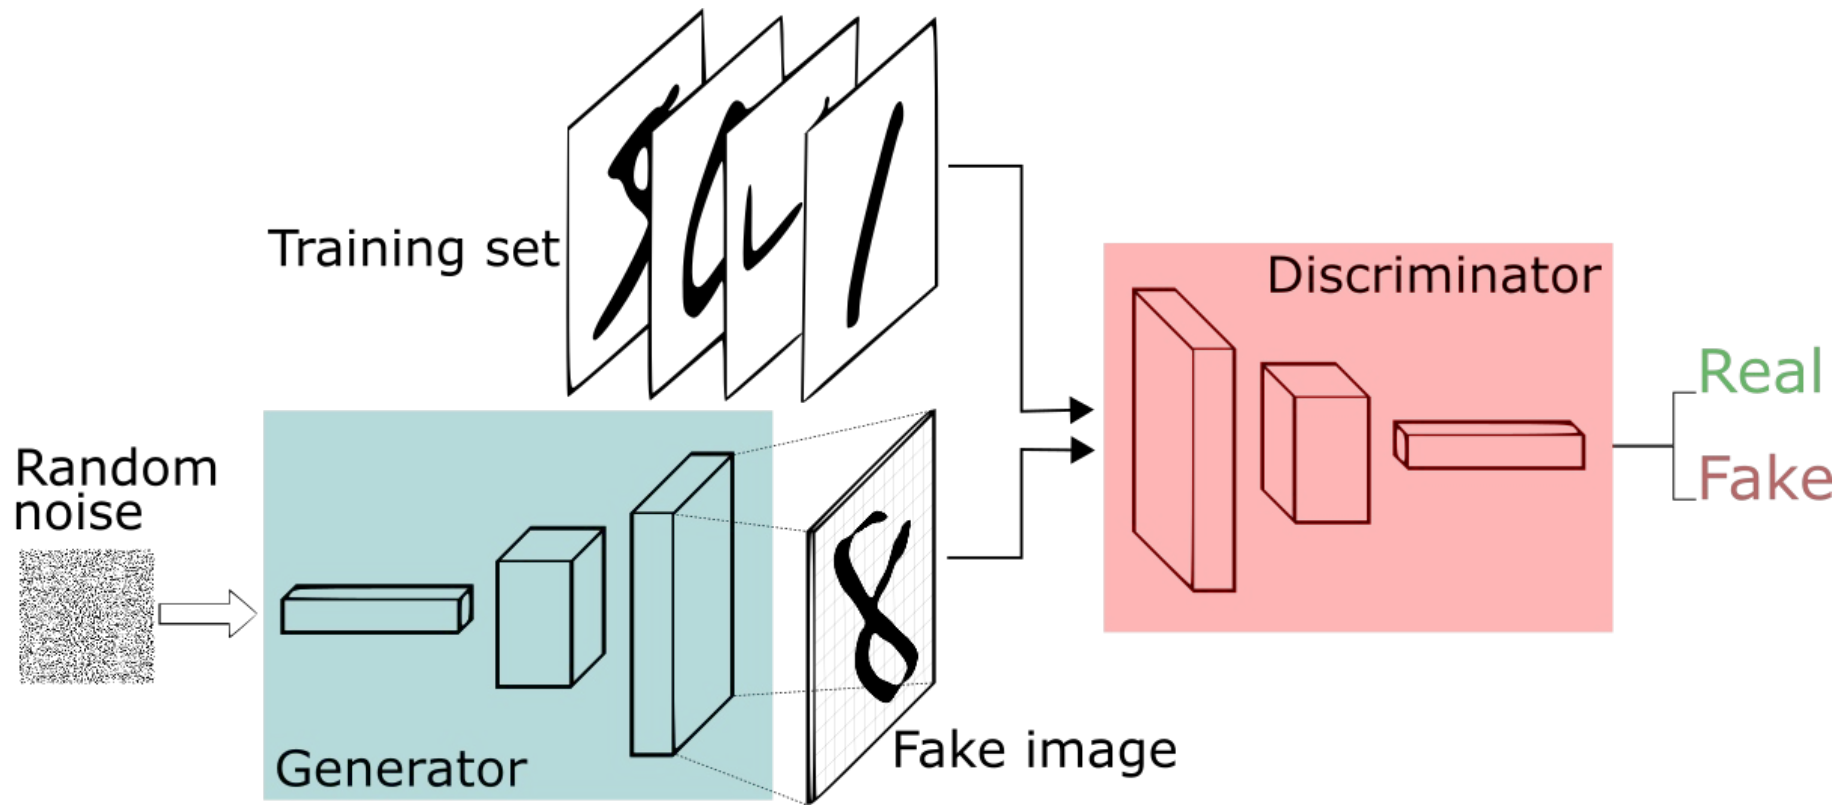
- Simple two-layer NN that assigns to categorical inputs a floating point vector
- Inputs that are closely associated will be close in the output space
- It can also find analogies


Autoencoders






Generative Adversarial Network (GAN)





Fantastic beasts and where to find them (2016)






Is this the real life? (2019)



NVIDIA, 2019

<https://thispersondoesnotexist.com/>



(Small Ad)

Study design consultation (free)
+ drop-in sessions every week @ all 6 sites

Project support (user fee 800 kr/h)


Long-term support (free; selection by peer review)

Partner projects (cost: part-time employment)

PhD advisory program (one call per year)



NBS
NATIONAL BIOINFORMATICS
INFRASTRUCTURE SWEDEN



ny metod som ers... x safex.se x Sprängmedel Dyn... x Sidan kunde inte... x stenspräckaren - S... x Stenspräckning oc... x NBIS - National Bi... x

nbis.se

Appar → Störningsinformation → NYheter - DN.SE → Latest Headlines → SJ Internet ombord → SJ → Lp → SL → SSS → Gmail → U → Cst → Lp → Google Maps → Övriga bokmärken

Support ▾ Infrastructure ▾ Training ▾ News ▾ About ▾

- General information
- Consultation
- Support
- Long-term Support
- Genome Assembly and Annotation
- Apply here!


NBIS

NATIONAL BIOINFORMATICS INFRASTRUCTURE SWEDEN

NBIS is a distributed national bioinformatics infrastructure, supporting life sciences in Sweden

nbis.se/support

Royex_gen2_brosc...pdf ^ Royex_cartridges_t...pdf ^ Visa alla x



Expertise

Genomics

Systems biology

Proteomics

Machine Learning

Multi-omics integration

Metabolomics

Phylogenomics

Biostatistics

Computational method development